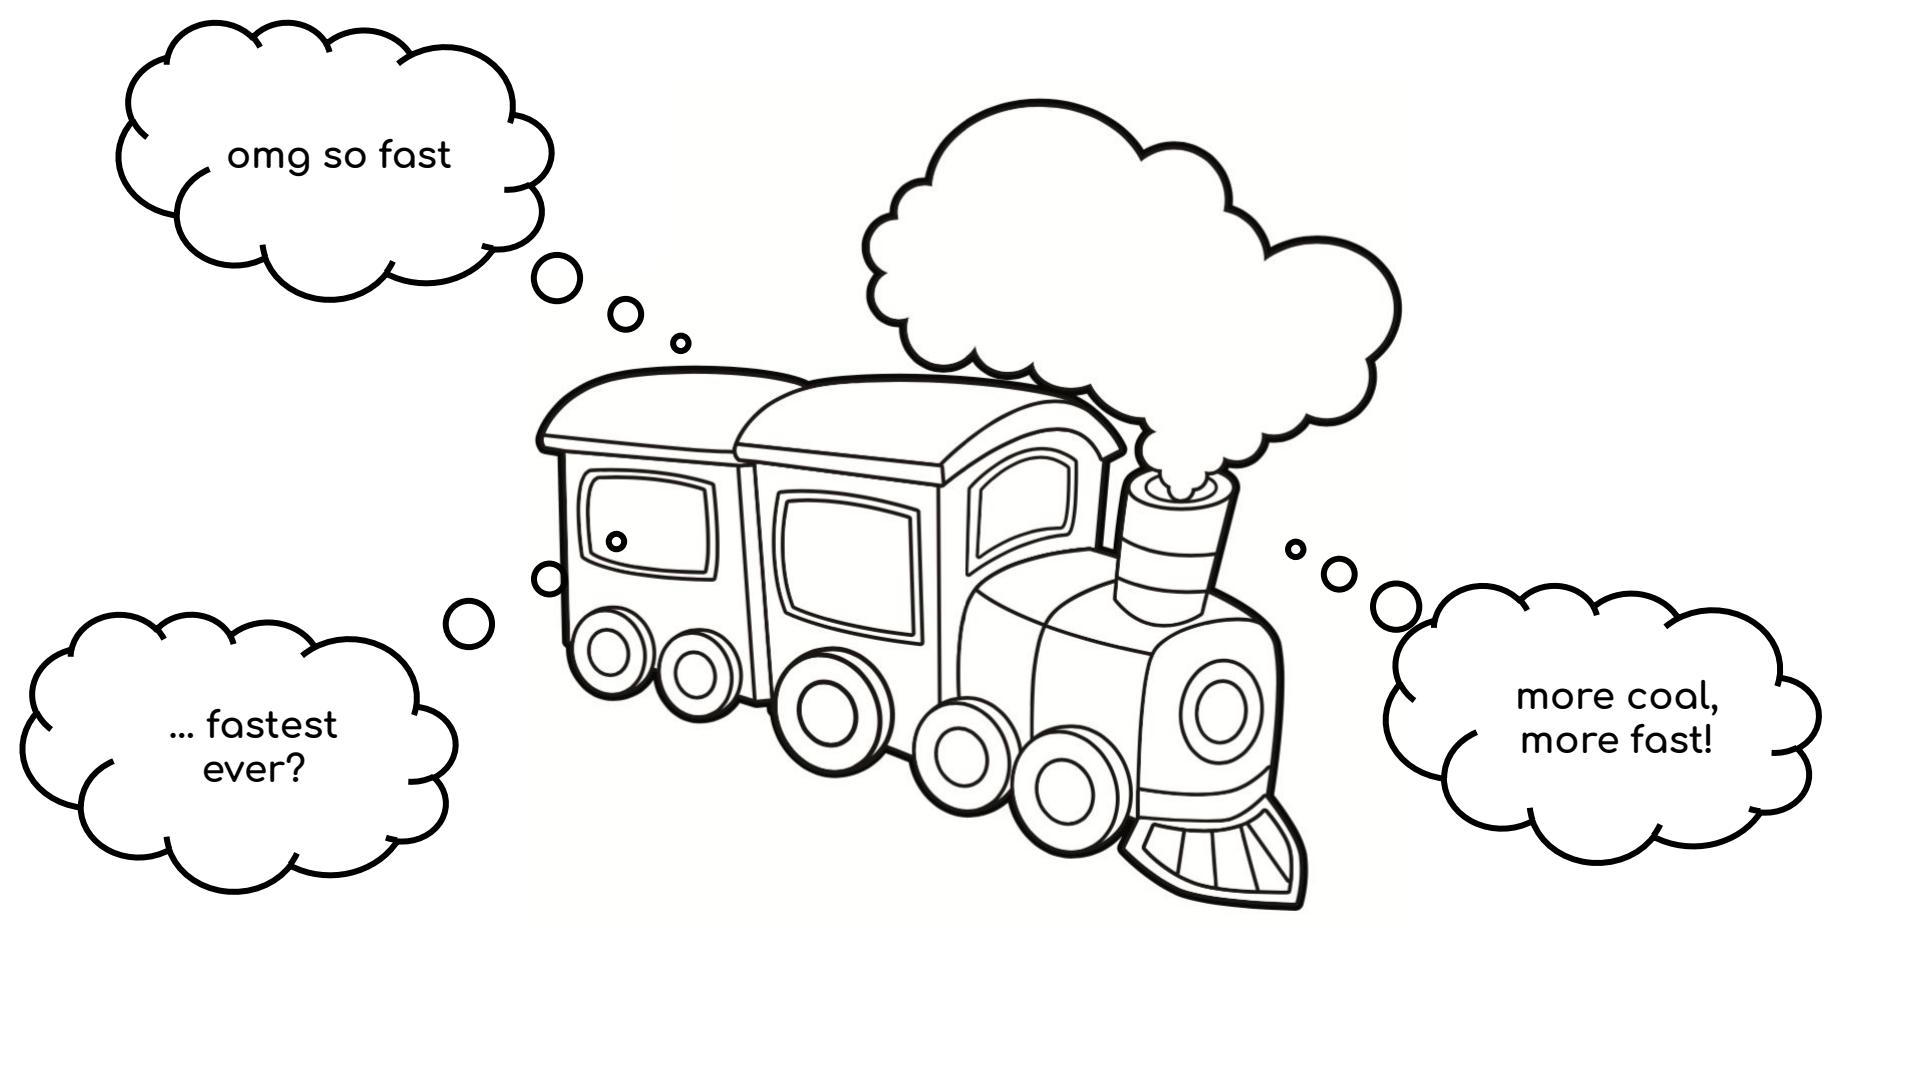


Machine learning can't do the thinking

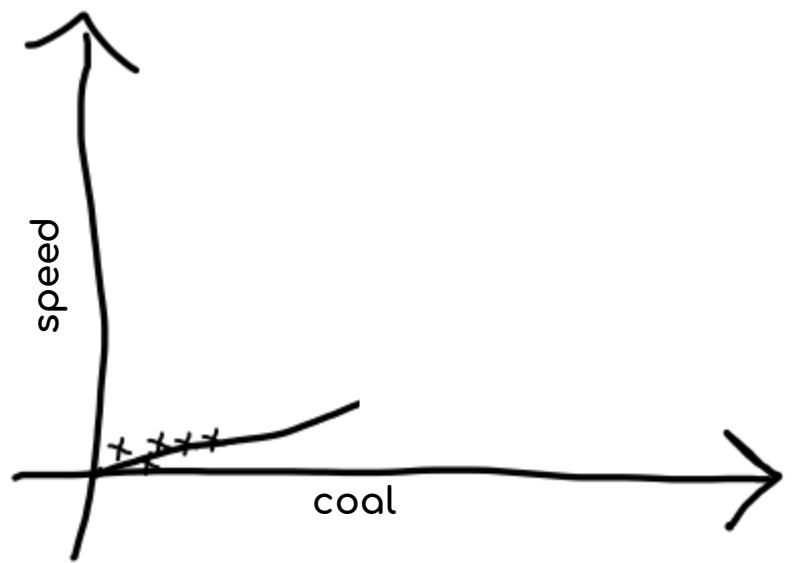
Inga Strümke, December 2019



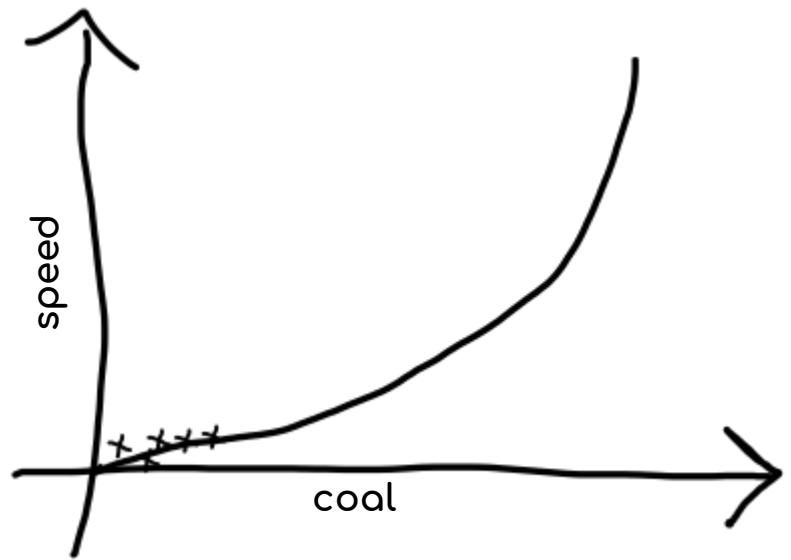
omg so fast

... fastest
ever?

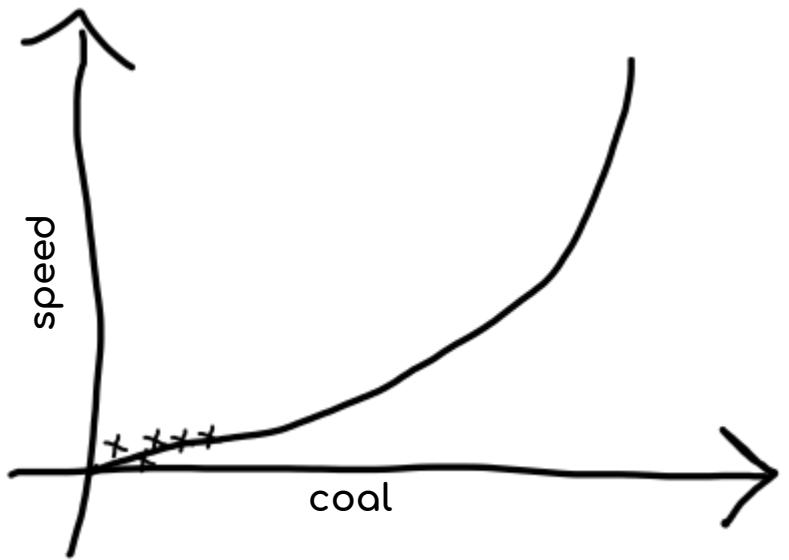
more coal,
more fast!



Q: will $v_{train} \rightarrow \infty$ if
the engine gets big
enough?



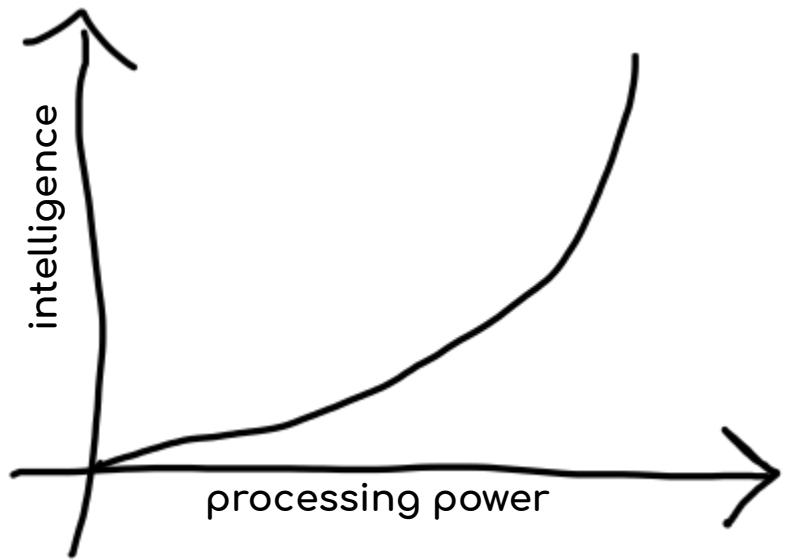
Q: will $v_{train} \rightarrow \infty$ if
the engine gets big
enough?



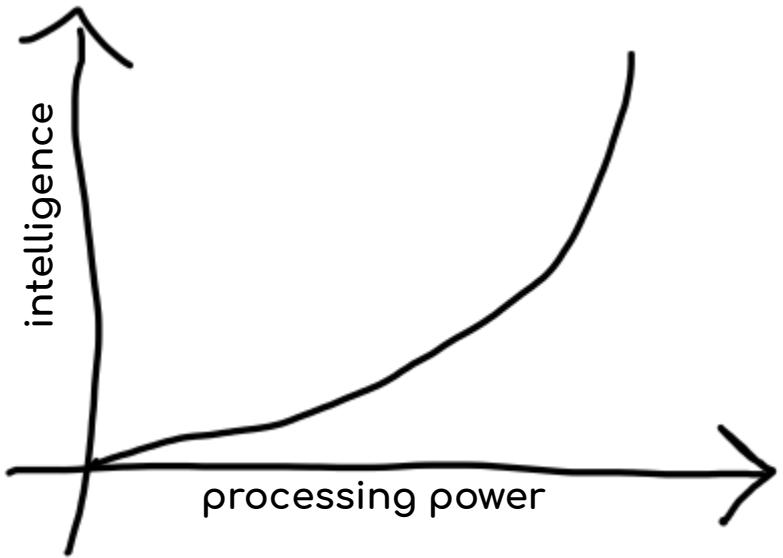
A: No, because

1. Architecture:
Rails, friction, air
resistance, locomotive
2. Physical limitations:
Sound barrier, light
speed,

Q: will AGI arise if
we increase
processing power?

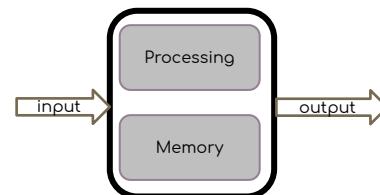


Q: will AGI arise if
we increase
processing power?



A: ... :)

1. Architecture



Moravec's paradox

2. Physical limitations?

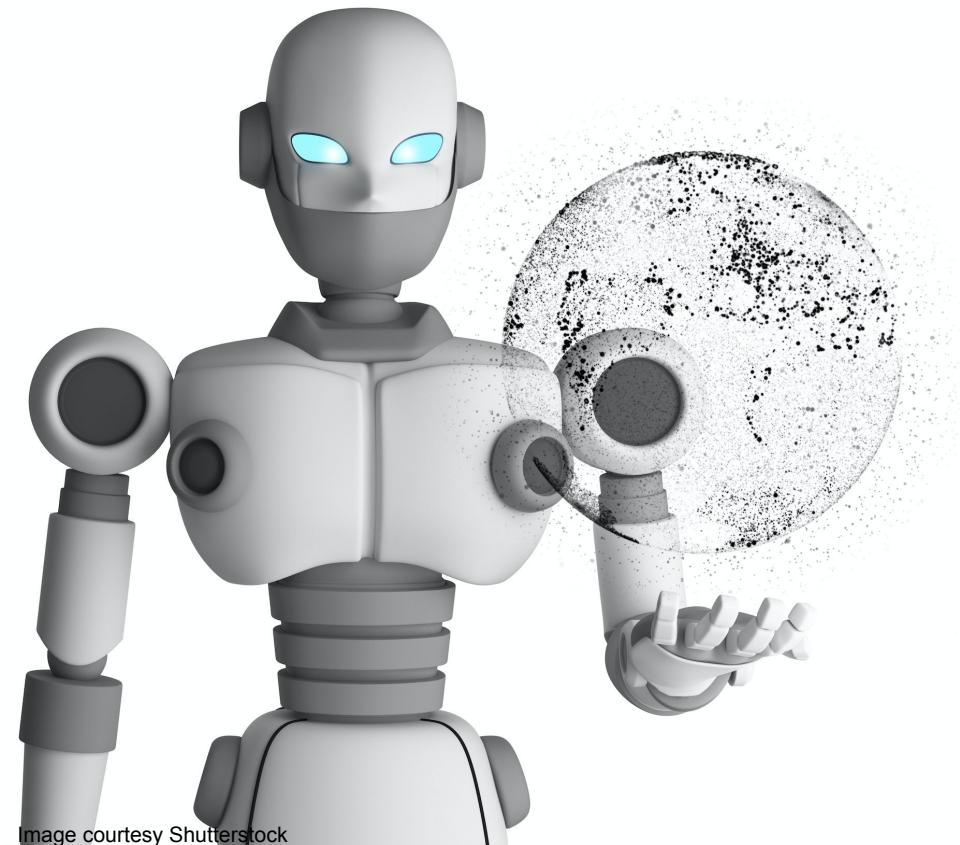
No need to worry about AGI yet (unfortunately)

How then to achieve AI world domination?

(and how to make this a good world)



The road to AI world domination:



- DATA

Luca Pacioli

The accountant who changed the world

Venice 14xx: Invented “double-entry bookkeeping”

E.g.: You sell 48 pepper for 3 gold coins ⇒

Gold coins	Pepper
10	200
13	152



Changed the world for two reasons:

1. It works **REALLY** well



Before, people kept diaries and counted their money at the end of the day

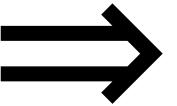
2. Part of a powerful system



The large and influential systems adopted it so it became a common standard



Use
standardised
format



Containers

Imagine a future where you say

“it used to take me a week to get all the data for my project. These days, it takes two hours”

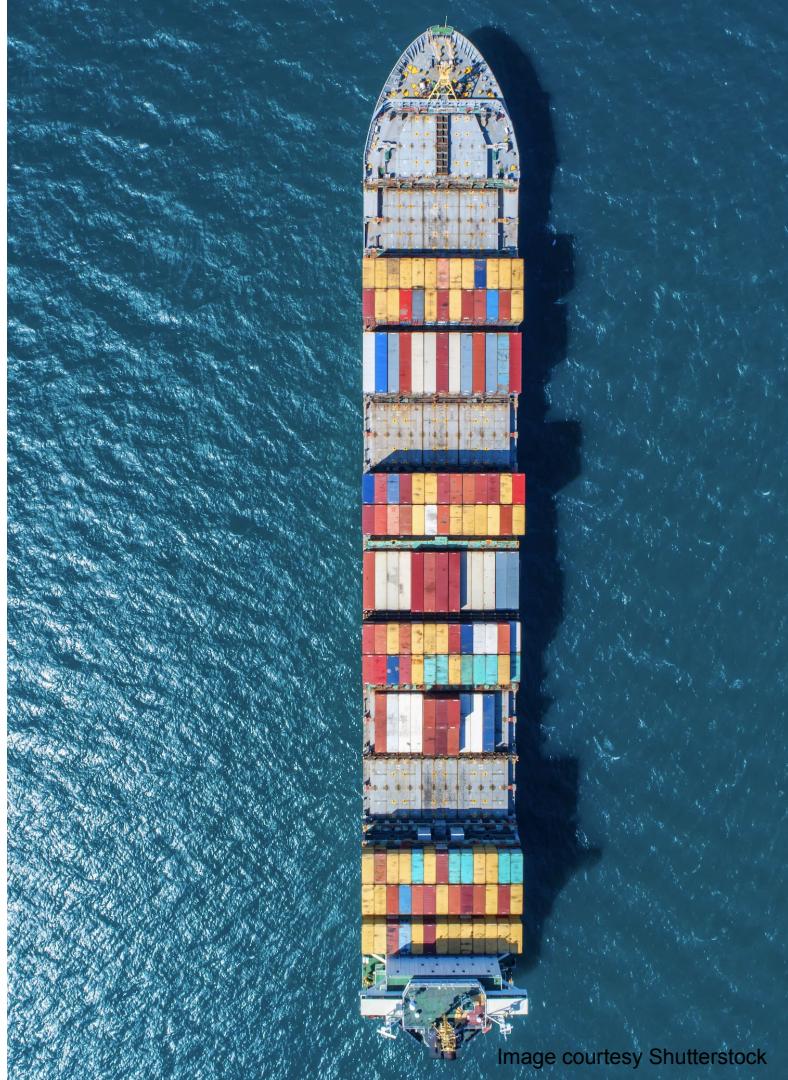


Image courtesy Shutterstock

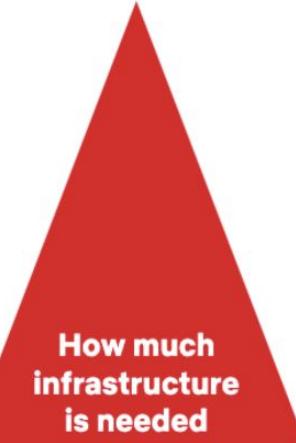
About two years ago, we, at our newly formed Machine Learning Infrastructure team started asking our data scientists a question: “What is the hardest thing for you as a data scientist at Netflix?” We were expecting to hear answers related to large-scale data and models, and maybe issues related to modern GPUs. Instead, we heard stories about projects where getting the first version to production took surprisingly long — mainly because of mundane reasons related to software engineering. We heard many stories about difficulties related to data access and basic data processing. We sat in meetings where data scientists discussed with their stakeholders how to best version different



Model Development
Feature Engineering
Model Operations
Versioning
Architecture
Job Scheduler
Compute Resources
Data Warehouse



Savin Goyal · 3rd
Machine Learning Infrastructure at Netflix



Once we've solved the data issues, what could ruin it for us?



1. Not using AI where we should use it
(missed use)
2. Using AI in a bad way
(misuse 1: being unethical)
3. Using AI the wrong way
(misuse 2: not understanding it)

Example from Norway Targeted marketing



Peppes Pizza:

Largest pizza restaurant chain in
Scandinavia



Example from Norway Targeted marketing



Peppes Pizza:

Caused a small media storm and would have been a **reputational disaster** for a smaller/younger company





Hong Kong protesters use lasers to confuse police and to interfere with their cameras. (Photos posted online)

HONG KONG / CCTV - 08/06/2019

Hong Kong protesters use lasers to confuse police and damage cameras





This is awesome.
We fight crime like
never before, and
society gains from
it. Everybody's
better off!



Get your hands off MY
data

Different ethics, different metrics, but one thing is certain:

The “eastern model” for sharing data provides perfect conditions for machine learning.

The road to AI world domination:



- DATA ✓
- ETHICS

Exclusive: Google cancels AI ethics board in response to outcry

The controversial panel lasted just a little over a week.

By Kelsey Piper | Apr 4, 2019, 7:00pm EDT

[f](#) [t](#) [SHARE](#)



Our favourite bad guy these days: BIAS
(the part where an algorithm learns from the past)

Amazon reportedly scraps internal AI recruiting tool that was biased against women



Bias is not as simple as over- and underrepresented groups

Task: Determine risk score for selecting which patients to treat first

Problem: White patients receive higher risk scores, than black patients who are equally sick

Why?

- Medical histories are used to predict **how much people will cost the health-care system**
- Black patients have historically incurred **lower costs** than white patients with the same conditions

⇒ The problem is not unequal representation in the data, but the **loss metric** in the cultural and historical (i.e. data) context

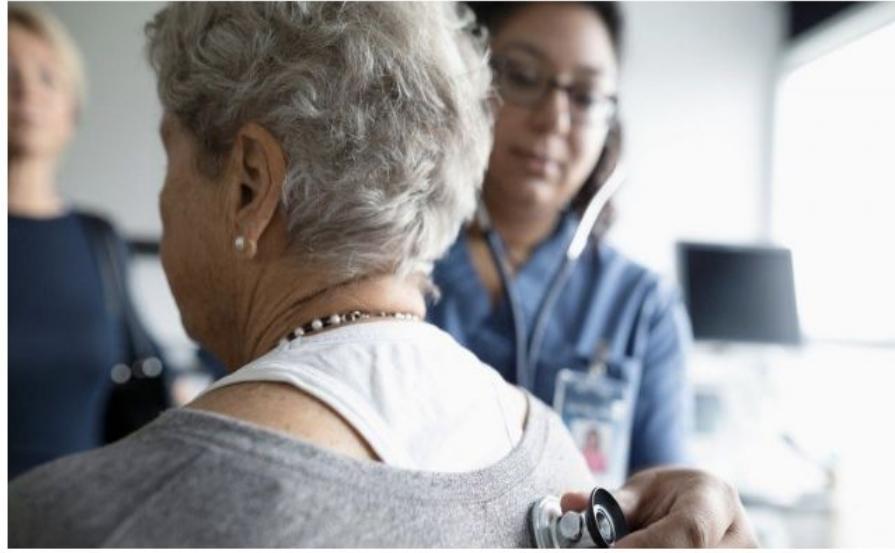
Easy fix: Disparity can be reduced >80% by predicting also the number of times a chronic condition will reoccur

Algorithmic bias can be corrected, **if you can catch it.**

MIT
Technology
Review

Artificial Intelligence Oct 25

A biased medical algorithm favored white people for health-care programs



Bias is hard to get rid of.

“ “

... when Sussman was a novice, Minsky once came to him as he sat hacking at the PDP-6.

"What are you doing?", asked Minsky.

"I am training a randomly wired neural net to play Tic-tac-toe", Sussman replied.

"Why is the net wired randomly?", asked Minsky.

"I do not want it to have any preconceptions of how to play", Sussman said.

Minsky then shut his eyes.

"Why do you close your eyes?" Sussman asked his teacher.

"So that the room will be empty."

At that moment, Sussman was enlightened.

What else is new?



Marvin Minsky, "Father of Artificial Intelligence"

ECONOMIC VIEW

Biased Algorithms Are Easier to Fix Than Biased People

Racial discrimination by algorithms or by people is harmful — but that's where the similarities end.



AI - in the sense of advanced data analysis - is actually good news

They highlight existing problems

The algorithms don't decide what is ethical

“

Bias vs fairness

While **bias** is a well-defined concept and can be controlled, **fairness** is a societal construct where different interests are put up against each other

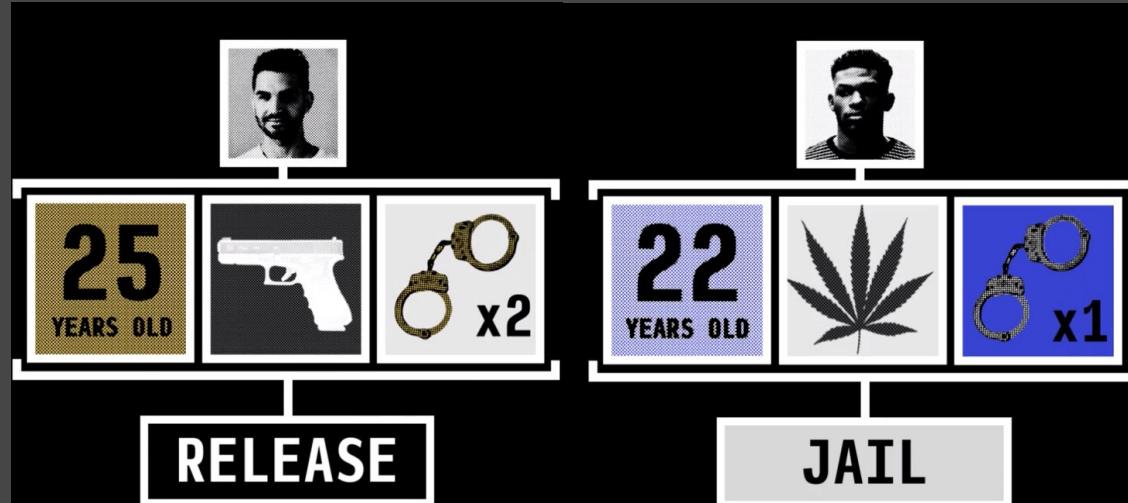
COMPAS

- **Task: Calculate risk score** → should defendant be kept in jail while awaiting trial?
- Aim: Make the criminal legal system fairer, by replacing the judges' intuition with a testable tool
- *Algorithm does not have access to protected attributes (ethnicity, age, gender, disability, etc)*

No ethnicity info in the training data

Can the algorithm still have racial bias?

Why?



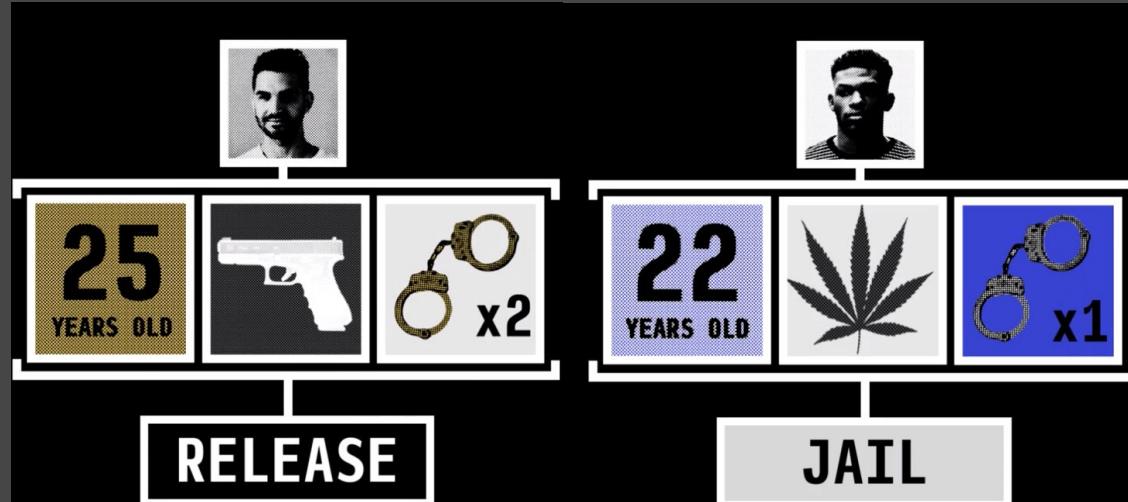
COMPAS

- **Task: Calculate risk score** → should defendant be kept in jail while awaiting trial?
- Aim: Make the criminal legal system fairer, by replacing the judges' intuition with a testable tool
- *Algorithm does not have access to protected attributes (ethnicity, age, gender, disability, etc)*

The algorithm has racial bias, as [demonstrated in 2016](#) by ProPublica.

Why?

Because there are predictors for ethnicity in the data set - i.e. the data could be used to build an ethnicity-predictor.



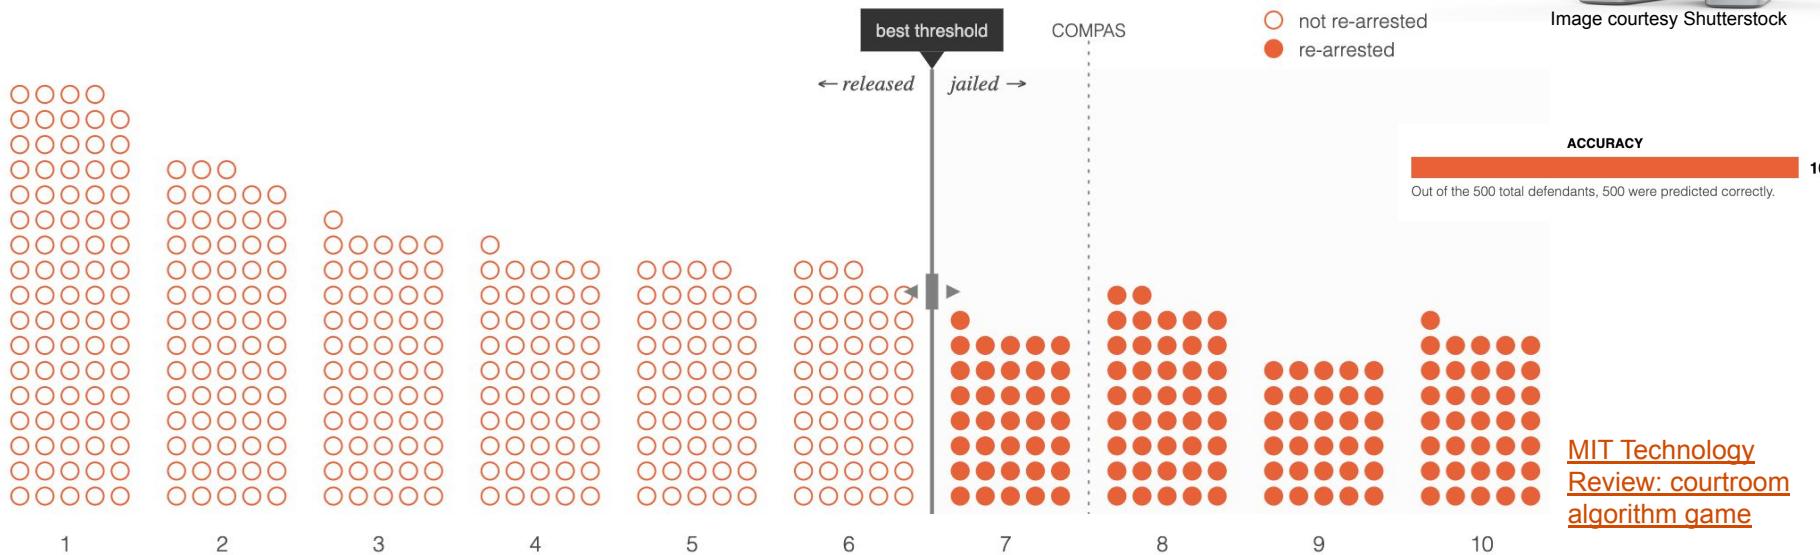
COMPAS - can we make it fair?

Definition

ERROR RATE: How many people were not jailed but still reoffended



Image courtesy Shutterstock



COMPAS - can we make it fair?

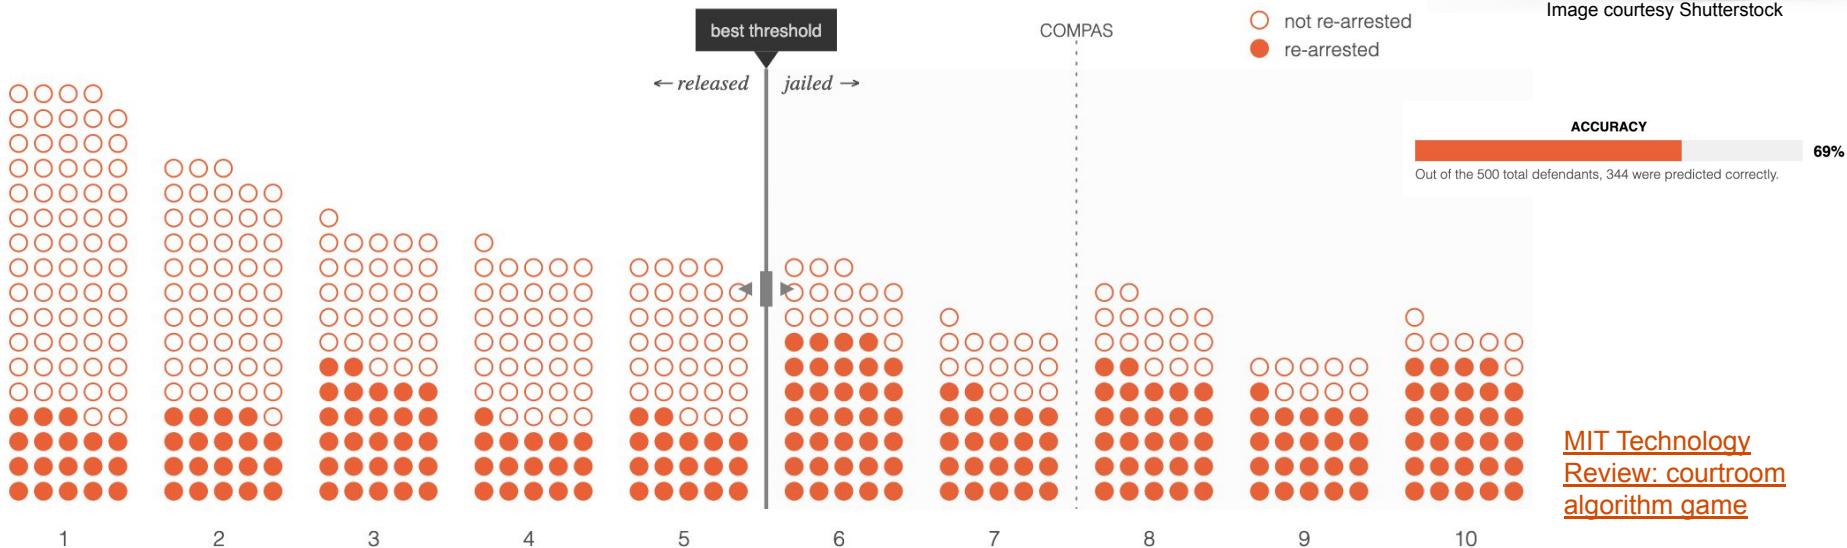
Definition

ERROR RATE: How many people were not jailed but still reoffended

Can't predict with 100% certainty how a person will act



Image courtesy Shutterstock



COMPAS - can we make it fair?

Definition

ERROR RATE: How many people were not jailed but still reoffended

Complication: White and black people are jailed with different frequencies. Reason? Historical discrimination?

Consequence: Two groups arise in the model. We wish this wasn't the case - but now that it is, we have to deal with it.

WHAT IS FAIR?

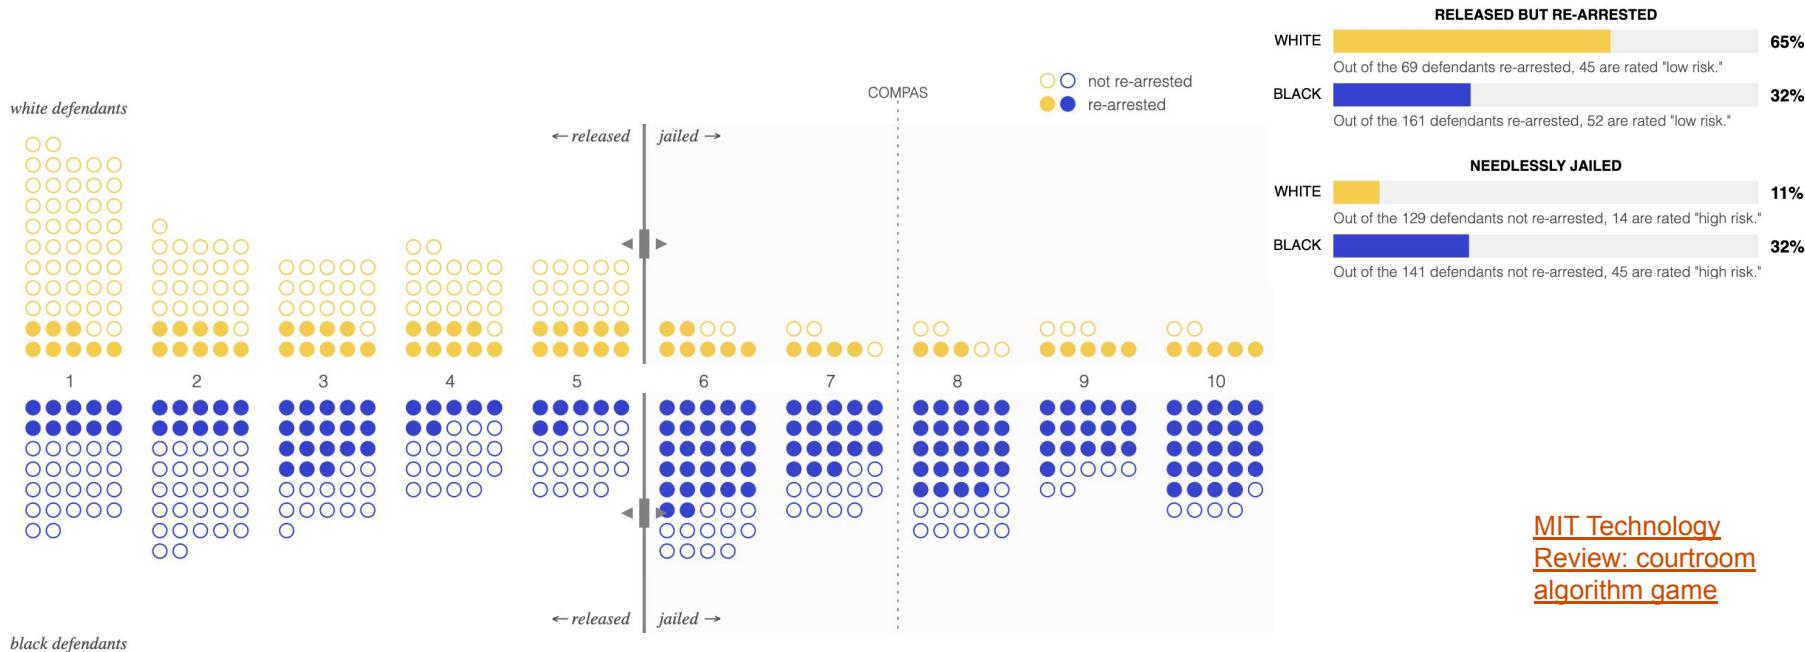
1. Keep the error rates comparable between groups (wrongfully arrest the same number of black and white people)
2. Treat people with the same risk scores in the same way

COMPAS - can we make it fair?

FAIRNESS:

1. Keep the error rates comparable between groups (wrongfully arrest the same number of black and white people)
2. Treat people with the same risk scores in the same way

PROBLEM: The error rates are different - more black than white people are needlessly arrested!

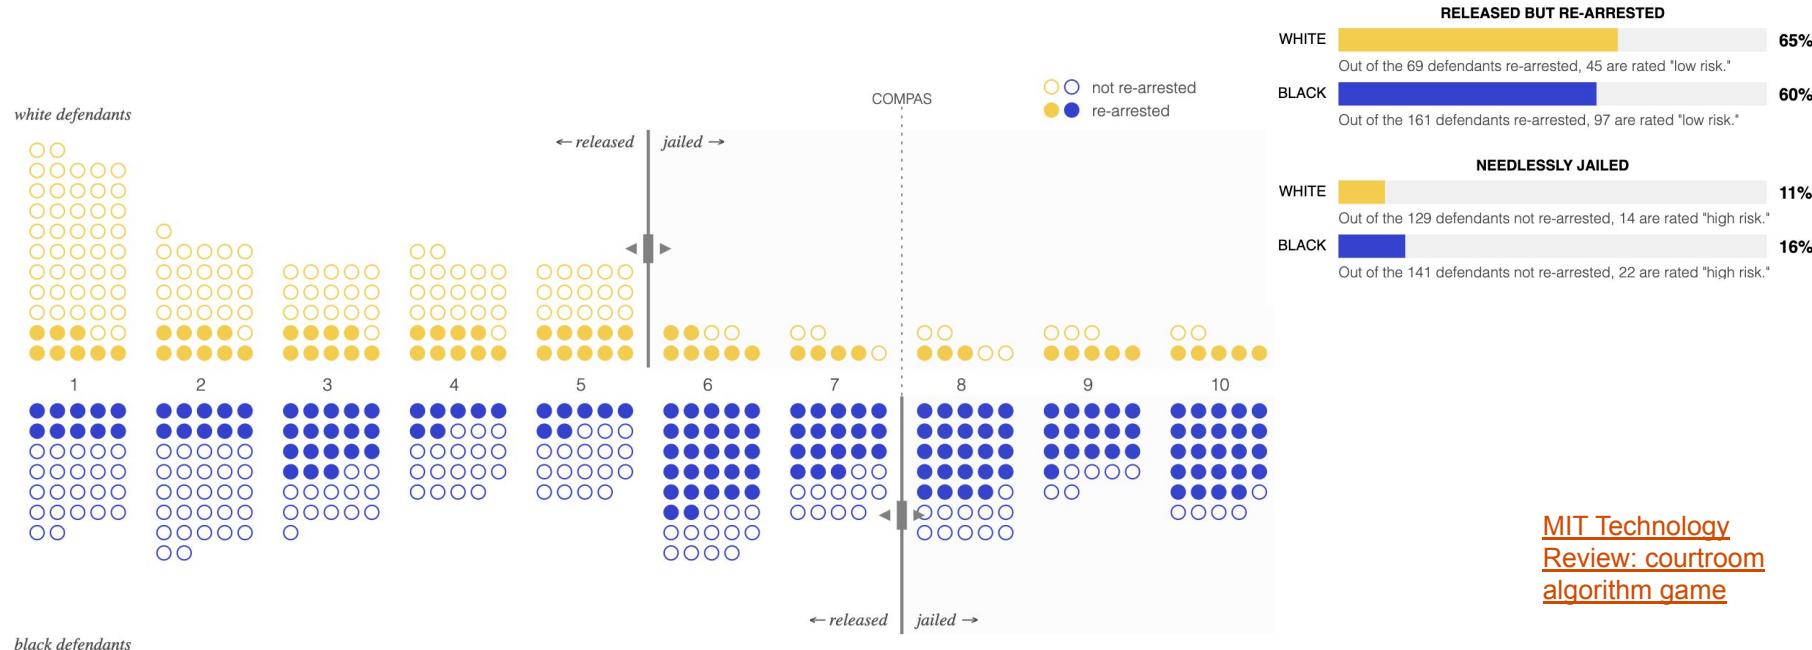


COMPAS - can we make it fair?

FAIRNESS:

1. Keep the error rates comparable between groups (wrongfully arrest the same number of black and white people)
2. Treat people with the same risk scores in the same way

PROBLEM: Two people with the same risk score are treated differently based on skin colour



There are at least 20 definitions of fairness, and when we choose one, we violate some aspect of the others.

Some fairness metrics examples:

Statistical Parity

Difference of the rate of favorable outcomes for the unprivileged group to the privileged group



Group fairness

Perspective: All groups have the same underlying abilities

Disparate Impact

Ratio of the rate of favorable outcome for the unprivileged group to the privileged group.

Equal Opportunity

Difference of true positive rates between the unprivileged and the privileged groups

Average Odds

Average difference of false positive rate and true positive rate between unprivileged and privileged groups.



Individual fairness

Perspective: The observations reflect ability

Fairness is a societal construct

It cannot be left to developers and programmers to decide which fairness definition to use or which group to protect

Answering such questions requires a democratic process



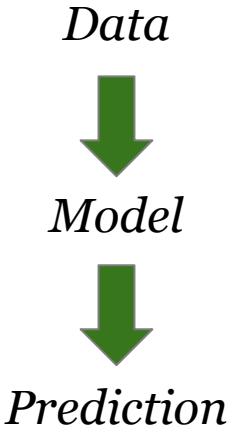
Finally: Not understanding



Image courtesy Shutterstock

Finally: Not understanding

Let's take a step back and think about what machine learning *does*



Machine learning is about predicting the future **based on the past**

... so it's strictly speaking association more than prediction



Image courtesy Shutterstock

“

**machines will always beat humans
in closed systems**

GARRY KASPAROV

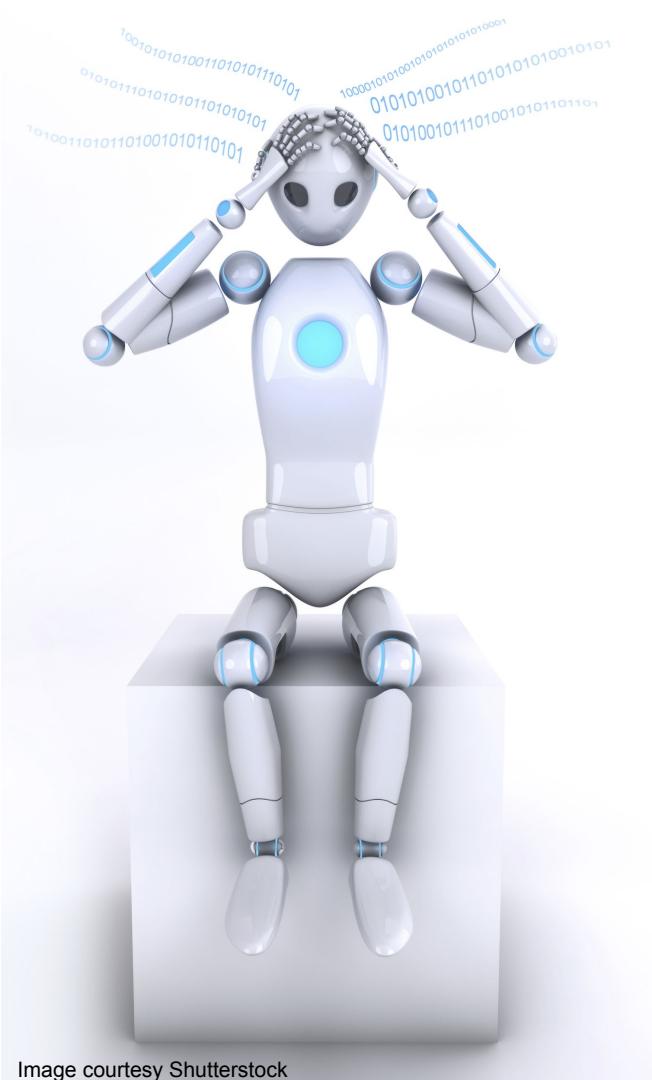
FORMER WORLD CHESS CHAMPION

Observations:

- Reality can be described
 - Machine learning can do this

Coincidence??

$\sin(x+y) = \sin x \cos y + \sin y \cos x$ $(\ln(x))' = x^{-1}$ $\sin x = 0,5$ $\int \frac{dx}{x^2 \pm a^2} = \ln|x \pm \sqrt{x^2 \pm a^2}| + C$ $(a+b)^c = a^c + 2ab$
 $(1+x)^a = 1 + \sum_{n=1}^{\infty} \binom{a}{n} \cdot x^n$ $\frac{a}{\sin A} = \frac{b}{\sin B}$ $e^{i\pi} + 1 = 0$ $\operatorname{tg} \alpha = \frac{\sin \alpha}{\cos \alpha}$ $\sin^2 \alpha + \cos^2 \alpha = 1$
 $\operatorname{Im} z = 3+2i$ $\operatorname{Re} z = \operatorname{C}_n^a = \frac{n!}{(n-a)!a!}$ $\left| \begin{array}{ccc} a & b & c \\ d & e & f \\ g & h & i \end{array} \right| = -\left| \begin{array}{ccc} a & b & c \\ d & e & f \\ g & h & i \end{array} \right| - \left| \begin{array}{ccc} a & b & c \\ d & e & f \\ g & h & i \end{array} \right|$ $\pi =$
 $x \in [3; +\infty)$ $\sin^2 \alpha + \cos^2 \alpha = 1$ $\sinh x = -i \sin(ix)$ $(e^x)' = e^x$ $(f(x))' = \lim_{\Delta x \rightarrow 0} \frac{\Delta f}{\Delta x}$
 $x \in (-\infty, -2)$ $(e^x)' = e^x$ $\lim_{x \rightarrow 0} \frac{\sin x}{x} = 1$ $f(x) = \frac{1}{\sigma \sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$ $\int_0^b xf'(x) dx$
 $a^2 = b^2 + c^2 - 2bc \cos A$ $\alpha = \operatorname{atan}(b/a)$ $\operatorname{tg} \alpha = \frac{b}{a}$ $(\sin x)' = \cos x$ $(C)' = 0$
 $y = \sin x$ $D = b^2 - 4ac$ $\left(\begin{array}{cc} a_1 & b_1 \\ a_2 & b_2 \end{array} \right) \cdot \left(\begin{array}{cc} c_1 \\ c_2 \end{array} \right) = \left(\begin{array}{c} a_1c_1 + b_1c_2 \\ a_2c_1 + b_2c_2 \end{array} \right)$ $i = \sqrt{-1}$ $\lim_{x \rightarrow 0} \frac{\sin x}{x} = 1$
 $\frac{1}{2^n} - 2$ $e^x = 1 + \sum_{n=1}^{\infty} \frac{x^n}{n!}$ $\int x^n dx = \frac{x^{n+1}}{n+1} + C$ $\sinh(x) = \frac{e^x - e^{-x}}{2}$ $\sin 90^\circ = 1$
 $\sin x = \operatorname{Im}\{e^{ix}\}$ $\operatorname{cosh}(x) = \frac{e^x + e^{-x}}{2}$ $A_n^k = \frac{n!}{(n-k)!}$ ∞ $\lim_{x \rightarrow 2} \frac{x^2 + 3x - 10}{x-2}$
 $X = 1$ $\log_{10} x = \frac{1}{P} \log_a x$ $a \perp m, a^{q(m)} \equiv 1 \pmod{m}$ $\log(ab) = \log a + \log b$ $S = 4\pi R^2$
 $X! = 1 \cdot 2 \cdots x$ $1 \quad 1 \quad 1 \quad 1$ $h = D \cdot \operatorname{tg} \alpha$ $V = \frac{4}{3}\pi R^3$ $\sqrt{2} = 1,41$
 $a \cap b = \emptyset$ $1 \quad 1 \quad 1 \quad 1$ $S = \frac{1}{2} ab \sin \alpha$ $(e^x)' = e^x$ \sum
 $1 \quad 1 \quad 1 \quad 1$ $1 \quad 4 \quad 6 \quad 4 \quad 1$ $y = x^2$ $\int_a^b f(x) dx$ $S_k = \sum_{i=1}^k a_i$
 $1 \quad 5 \quad 10 \quad 15 \quad 5 \quad 1$ $1 \quad 6 \quad 15 \quad 20 \quad 15 \quad 6 \quad 1$ $\cos 2\alpha = 2 \cos \alpha - 1$ $y = |x-2|$ $\sin A = \frac{a}{c}$
 $\operatorname{tg} \alpha = \frac{\sin \alpha}{\cos \alpha}$ $\sqrt[n]{x_1 x_2 \cdots x_n} \leq \frac{x_1 + x_2 + \cdots + x_n}{n}$ $\sum_{n=0}^k \frac{f^{(n)}(a)}{n!} (x-a)^n$ $\cos A = \frac{b}{c}$
 $\cos(x+y) = \cos x \cos y - \sin x \sin y$

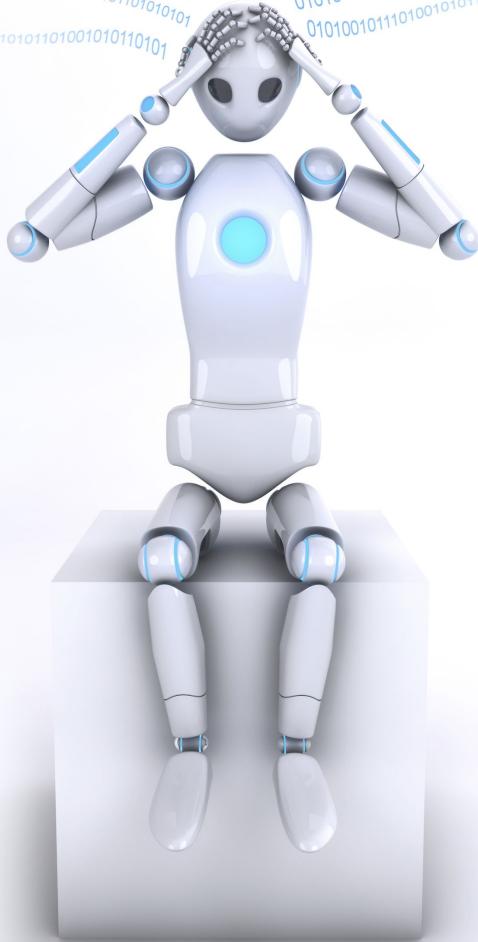


Machine learning is
correlation detection.

We still have to do the thinking

... and there is a huge effort to be
done also in XAI

10010101010011010101110101
100001010100101010101010001
01010100101101010101010101
0101001011101001010110101
101001101011101001010110101



Thank you!

inga.struemke@pwc.com